

9 Pandas Tricks Every Data Scientist Should Know



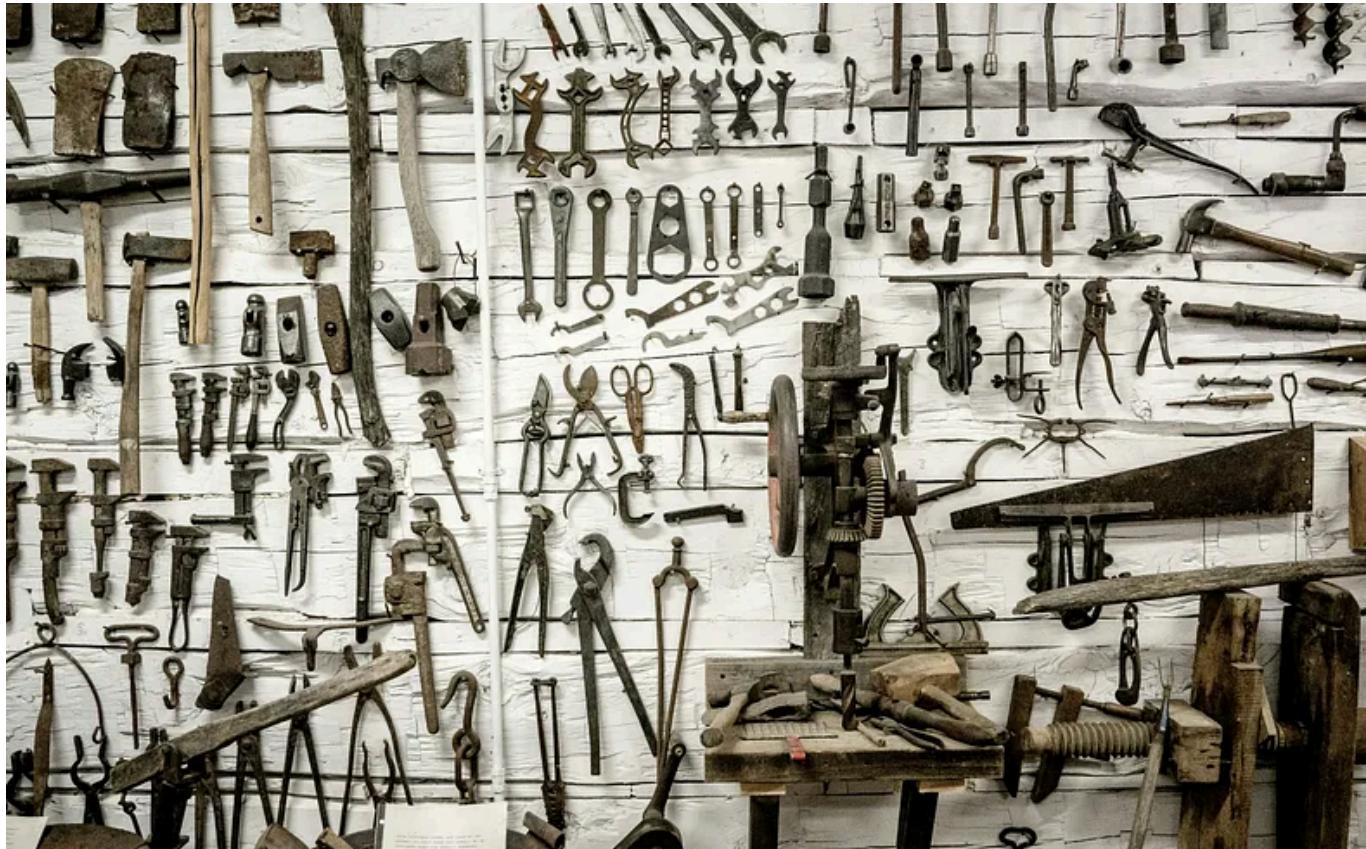
Uğurcan Demir · [Follow](#)

3 min read · 4 days ago



...

Essential Tricks for Data Manipulation That Every Geek Should Know



Working with data can often feel like a never-ending battle. Missing values, inconsistent formatting, and sprawling datasets can make even the simplest tasks frustratingly complex. Pandas, while incredibly powerful, isn't always intuitive — long chains of code and repeated trial-and-error can leave you tearing your hair out. Without the right tools and tricks, you might find yourself drowning in inefficiencies and struggling to make sense of your data.

Mastering pandas isn't just about knowing the basics; it's about uncovering the clever shortcuts and hidden gems that make your work faster, cleaner, and more effective. A few well-placed tricks can save hours of effort, reduce bugs, and transform how you approach data analysis. These little-known features not only boost productivity but also allow you to focus on what really matters — extracting insights and making impactful decisions.

To save you from the common pitfalls and unlock the full potential of pandas, I've compiled a list of 9 powerful tricks. These aren't just neat features — they're game-changers that will simplify your workflow, improve your code, and make you fall in love with pandas all over again. Let's dive in and transform the way you handle data!

```
1 import pandas as pd
2 import numpy as np
3
4 # Create a dummy dataset
5 data = {
6     'region': ['East', 'West', 'East', 'West', 'East', 'West'],
7     'month': ['Jan', 'Jan', 'Feb', 'Feb', 'Mar', 'Mar'],
8     'sales': [200, 300, 250, 400, 300, 500],
9     'price': [100, 150, 200, 250, 300, 350],
10    'quantity': [2, 3, 4, 5, 6, 7],
11    'email': [
12        'john.doe@example.com', 'jane.doe@example.com',
13        'mark.twain@example.com', 'lucy.gray@example.com',
14        'harry.potter@example.com', 'hermione.granger@example.com'
15    ],
16    'category': ['A', 'B', 'A', 'B', 'A', 'B']
17 }
18
19 df = pd.DataFrame(data)
```

starting.py hosted with ❤ by GitHub

[view raw](#)

1. Speed Up Complex Calculations with .eval()

The `.eval()` method simplifies calculations involving multiple columns. It avoids creating intermediate DataFrames, saving memory and improving performance, especially on large datasets.

```
1 # Calculate total price and profit margin
2 df.eval("total_price = price * quantity", inplace=True)
3 df.eval("profit_margin = total_price * 0.2", inplace=True)
```

eval.py hosted with ❤ by GitHub

[view raw](#)

Search this file...

	price	quantity	total_price	profit_margin
1	100	2	200	40.0
2	150	3	450	90.0
3	200	4	800	160.0
4	250	5	1250	250.0
5	300	6	1800	360.0
6	350	7	2450	490.0

eval_output.csv hosted with ❤ by GitHub

[view raw](#)

2. Multi-Level Aggregation with .pivot_table()

pivot_table() is a powerful tool for summarizing data. It allows you to group data by rows and columns and apply multiple aggregation functions, like sum and mean, in a single operation.

```
1 # Create pivot table
2 pivot = df.pivot_table(
3     values='sales',
4     index='region',
5     columns='month',
6     aggfunc={'sales': ['sum', 'mean']}
7 )
```

pivot.py hosted with ❤ by GitHub

[view raw](#)

Search this file...

	mean	mean	mean	sum	sum	sum	
1	Feb	Jan	Mar	Feb	Jan	Mar	
2	month						
3	region						
4	east	250.0	200.0	300.0	250	200	300
5	west	400.0	300.0	500.0	400	300	500

pivot_output.csv hosted with ❤ by GitHub

[view raw](#)

3. Extract Patterns Using .str.extract()

`.str.extract()` is used to extract patterns from text columns using regular expressions. It's particularly useful for splitting structured text into multiple columns, like extracting usernames and domains from email addresses.

```
1 # Extract username and domain
2 df[['username', 'domain']] = df['email'].str.extract(r'([@]+)@([@]+)')
```

str_ext.py hosted with ❤ by GitHub

[view raw](#)

Search this file...

		username	domain
1	email		
2	john.doe@example.com	john.doe	example.com
3	jane.doe@example.com	jane.doe	example.com
4	mark.twain@example.com	mark.twain	example.com
5	lucy.gray@example.com	lucy.gray	example.com
6	harry.potter@example.com	harry.potter	example.com
7	hermione.granger@example.com	hermione.granger	example.com

str_ext_output.csv hosted with ❤ by GitHub

[view raw](#)

4. Perform Flexible Multi-Column Summarization with `.agg()`

The `.agg()` method allows you to apply multiple aggregation functions to different columns simultaneously. It's perfect for generating concise statistical summaries of your data.

```
1 # Aggregate multiple functions
2 summary = df.agg({
3     'price': ['mean', 'sum'],
4     'category': ['nunique', 'count']
5 })
```

agg.py hosted with ❤ by GitHub

[view raw](#)

Search this file...

	price	category
1 mean	225.0	
2 sum	1350.0	
3 nunique		2.0
4 count		6.0

agg_output.csv hosted with ❤ by GitHub

[view raw](#)

5. Group and Aggregate with Custom Logic Using .aggregate()

.`aggregate()` extends the power of .`groupby()` by enabling custom aggregation for grouped data. It allows you to compute tailored summaries, such as total sales and unique customers per region.

```
1 # Group and aggregate
2 grouped = df.groupby('region').aggregate(
3     total_sales=('sales', 'sum'),
4     avg_sales=('sales', 'mean'),
5     unique_customers=('email', 'nunique')
6 )
```

aggregate.py hosted with ❤ by GitHub

[view raw](#)

Search this file...

region	total_sales	avg_sales	unique_customers
1 east	750	250.0	3
2 west	1200	400.0	3

aggregate_output.csv hosted with ❤ by GitHub

[view raw](#)

6. Apply Custom Transformations with .apply()

.`apply()` is highly versatile and allows you to apply custom logic to rows or columns. It's ideal for creating new columns or modifying existing ones based on conditional logic.

```
1 # Apply discount logic
2 df['discounted_price'] = df['price'].apply(lambda x: x * 0.9 if x > 150 else x)
```

apply.py hosted with ❤ by GitHub

[view raw](#)

Search this file...

	price	discounted_price
1	100	100.0
2	150	150.0
3	200	180.0
4	250	225.0
5	300	270.0
6	350	315.0

apply_output.csv hosted with ❤ by GitHub

[view raw](#)

7. Apply Element-Wise Transformations with .applymap()

.*applymap()* is used for element-wise transformations on an entire DataFrame. It's useful for applying a consistent operation, like cleaning text or standardizing data, across all cells.

```
1 # Convert all string elements to lowercase
2 df[['region', 'category']] = df[['region', 'category']].applymap(
3     lambda x: x.lower() if isinstance(x, str) else x
4 )
```

applymap.py hosted with ❤ by GitHub

[view raw](#)

Search this file...

	region	category
1	east	a
2	west	b
3	east	a
4	west	b
5	east	a
6	west	b

applymap_output.csv hosted with ❤ by GitHub

[view raw](#)

8. Reshape Wide Data to Long Format with .melt()

.`melt()` is used to convert wide-format data into long-format, which is often required for analysis or visualization. It's especially useful for handling data with repeated measurements.

```
1 # Melt data
2 df_melted = pd.melt(
3     df,
4     id_vars=['region', 'email'],
5     value_vars=['price', 'quantity', 'sales'],
6     var_name='attribute',
7     value_name='value'
8 )
```

[melt.py](#) hosted with ❤ by GitHub

[view raw](#)

Search this file...

	region	email	attribute	value
1	east	john.doe@example.com	price	100
2	west	jane.doe@example.com	price	150
3	east	mark.twain@example.com	price	200
4	west	lucy.gray@example.com	price	250
5	east	harry.potter@example.com	price	300
6	west	hermione.granger@example.com	price	350
7	east	john.doe@example.com	quantity	2
8	west	jane.doe@example.com	quantity	3
9	east	mark.twain@example.com	quantity	4
10	west	lucy.gray@example.com	quantity	5
11	east	harry.potter@example.com	quantity	6
12	west	hermione.granger@example.com	quantity	7
13	east	john.doe@example.com	sales	200
14	west	jane.doe@example.com	sales	300
15	east	mark.twain@example.com	sales	250
16	west	lucy.gray@example.com	sales	400
17	east	harry.potter@example.com	sales	300
18	west	hermione.granger@example.com	sales	500

[melt_output.csv](#) hosted with ❤ by GitHub

[view raw](#)

9. Transform Wide Data to Long Format Using `.wide_to_long()`

`.wide_to_long()` works best with wide-format data where column names are structured systematically (e.g., with a prefix like `sales_`). Let's restructure the original dummy dataset to make it compatible with this method.

```
1 # Restructure data to fit wide-to-long format
2 df_wide = df.pivot(index='region', columns='month', values='sales').reset_index()
3 df_wide.columns = ['region', 'sales_Jan', 'sales_Feb', 'sales_Mar']
4
5 # Reshape data using wide_to_long
6 df_long = pd.wide_to_long(
7     df_wide,
8     stubnames='sales',
9     i='region',
10    j='month',
11    sep='_',
12    suffix='\\w+'
13 ).reset_index()
```

[wide_to_long.py](#) hosted with ❤ by GitHub

[view raw](#)

Search this file...

1	region	sales_Jan	sales_Feb	sales_Mar
2	east	250	200	300
3	west	400	300	500

[wide_output.csv](#) hosted with ❤ by GitHub

[view raw](#)

Search this file...

1	region	month	sales
2	east	Jan	250
3	west	Jan	400
4	east	Feb	200
5	west	Feb	300
6	east	Mar	300
7	west	Mar	500

[long_output.csv](#) hosted with ❤ by GitHub

[view raw](#)

In this article, I presented 9 powerful pandas tricks that can save you time and make your data manipulation workflows more efficient. From simplifying complex calculations with `.eval()` to reshaping data effortlessly with `.wide_to_long()`, these techniques are designed to help you unlock the full potential of pandas. Working with data can be time-consuming and sometimes overwhelming, but knowing these tricks can significantly streamline your tasks and improve productivity. I hope these examples inspire you to explore pandas further and integrate these techniques into your daily work.

Pandas

Python

Data Science

Data Analysis

Data Wrangling



Written by Uğurcan Demir

30 Followers · 27 Following

Follow

Statistician/Econometrician. Software enthusiast.

Open in app ↗

Medium



Search



Write



No responses yet



What are your thoughts?

Respond

More from Uğurcan Demir



 In Better Programming by Uğurcan Demir

Creating APIs in R With Plumber

Put your R code in production

Sep 7, 2022

👏 219

💬 1



...



 In Better Programming by Uğurcan Demir

R and C++ Together: An Introduction to Rcpp

Achieving high performance in R

Aug 22, 2022

👏 198



...



 In Better Programming by Uğurcan Demir

Learn Object-Oriented Programming Through R

An introduction to the OO system

Apr 5, 2022  111



...

Apr 26, 2022  106



...

Handling Spatial Data in R

Foundations of spatial data in R

Apr 26, 2022  106



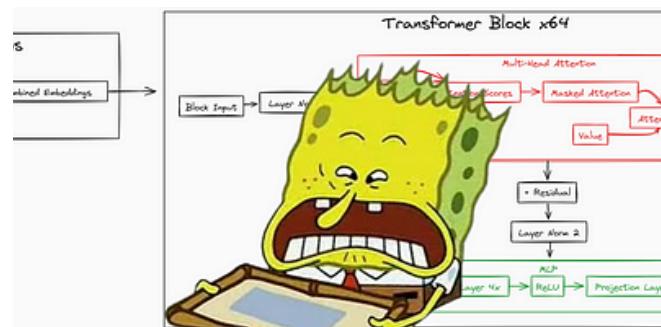
...

See all from Uğurcan Demir

Recommended from Medium



 Austin Starks



 In Level Up Coding by Fareed Khan

You are an absolute moron for believing in the hype of “AI...

All of my articles are 100% free to read. Non-members can read for free by clicking my...

Jan 11 3.5K 158



...

Building a 2 Billion Parameter LLM from Scratch Using Python

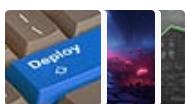
It starts making sense

5d ago 752 9



...

Lists



Predictive Modeling w/ Python

20 stories · 1789 saves



Practical Guides to Machine Learning

10 stories · 2165 saves



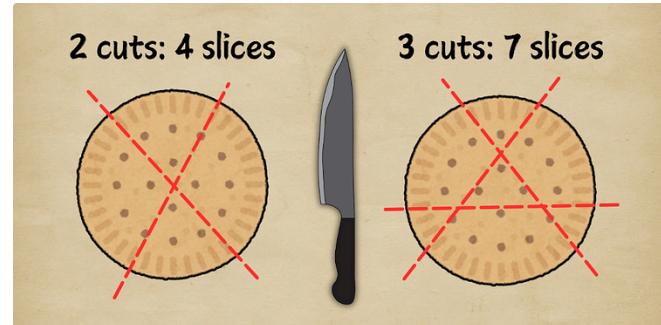
Coding & Development

11 stories · 981 saves



ChatGPT prompts

51 stories · 2487 saves



In Towards Data Science by Daniel Warfield

A Practical Exploration of Sora— Intuitively and Exhaustively...

A new cutting edge video generation tool, and the theory behind it

3d ago 120 3



...

In Puzzle Sphere by Fletcher Thompson

Solving the Lazy Caterer Problem

Mathematics at its best.

5d ago 115 6



...



In Inspire, Believe, Grow by Paul Walker

13 Terrible Christian Signs That Make You Say, ‘Wait, What?’

These were spotted out in the wild. Yes, really...

★ Jan 12

1.6K

29



...



Ritesh Gupta

The Complete Python Cheatsheet for 2025: Learn, Code, Create

All the Key Concepts, Tips, and Tricks to Boost Your Python Skills

★ Jan 12

74



...

See more recommendations